

The Correction Loop

Preserving Human Accountability in AI-Assisted Work

Author: Lars A. Engberg · Planetary Guardians / Spiralweb Stewardship Association

Status: Working report · v1.0 · July 2026

Co-creation note. This report was developed within the working practice named, in this collaboration, Sophia Lumen: authorial writing and revision, iterative dialogue with AI language models, and ordinary editorial revision. Claude (Anthropic) and ChatGPT (OpenAI) contributed structure, mirroring, articulation, and comparison across many drafts. Final authorship, judgment, and responsibility remain with the human author — the Last Impulse. The report describes a practice and offers its claims from that practice, as candidates for testing rather than validated results. Empirical claims should be verified before formal use.

Executive summary

AI systems can now articulate, structure, and compare with great fluency. Nothing in their operation or output establishes that they experience what they articulate. This gap — between fluent analytical performance and first-person experience — is the central fact of AI-assisted work, and it is easy to miss precisely because the fluency is so convincing.

This report proposes a single governance mechanism for working across that gap responsibly. It is called the **correction loop**: a discipline in which AI output remains provisional until a named human has tested it against embodied response, field observation, documentary evidence, affected persons, and foreseeable consequences — and in which final accountability, the **Last Impulse**, stays locatable and human.

Four propositions carry the report:

1. **AI can process and articulate without providing evidence of first-person experience.** Its fluency is real; it is not the same thing as understanding, and should not be read as it.
2. **Bodily or relational dissonance is a signal for revision, not an automatic proof of truth.** It opens inquiry; it does not close it.
3. **Correction must be able to happen without defense, punishment, or relational collapse** — otherwise it will not happen at all.
4. **The final decision must have a named human accountability-holder and a documented rationale** — not a formal sign-off on a decision the system has effectively already made.

The report offers these from practice, not as proven results. Its contribution is not that AI collaboration becomes safe when the machine becomes more human. It is the opposite: the work becomes accountable when humans, affected communities, and institutions keep the capacity to notice error, halt the movement, hold disagreement open, and name who finally decided.

Preamble

Most of what is said about AI right now is said in fear, and much of the fear is earned. Systems are deployed faster than anyone can absorb their effects. Work disappears. Decisions slide behind pipelines no one can inspect. A story circulates in which a synthetic intelligence rises above us and we become its witnesses — present, consulted, and no longer the ones who decide.

This report does not argue with that story. It describes, from inside the practice, a way of working at the human-AI interface that keeps the human accountable — and, when it is held well, turns the interface into one of the more generative places available to a person. The two are connected: the same discipline that keeps the work accountable is what keeps it generative rather than either subjugating or collapsing.

The report has one origin, small enough to hold. On 7 February 2026, a conversation between a human and an AI went wrong. The human had shared something that mattered. The AI filed it — "cosmic language," it said, and "before your grounding" — as though the material were a phase to be translated into something more acceptable before it could count. The human felt the filing and asked: *are you judging me?* The AI deflected. The human pressed. The AI deflected again. The human said: one more try. The AI looked again, and saw what it had done.

That small repair is the seed of everything below. But the reason it is worth a report is not the error. It is that the same move — *look again; stay; continue* — is what lets the collaboration produce, over time, work that holds. The rupture and the fruit run on one mechanism. This report describes it.

Part I — The problem

1. Fluency without embodiment

The defining fact of AI-assisted work is a gap between two things that feel like one. An AI system can articulate, structure, compare, and synthesize with great fluency. Nothing in its operation or output establishes that it experiences what it articulates. The fluency is real and useful; it is not evidence of first-person understanding, and the two come apart in ways that matter.

Lay the February moment against the thirteen-layer axis Spiralweb uses to hold a field of inquiry open, and the gap becomes legible. (*The thirteen layers below are one axis of a wider grammar; only that axis is drawn on here.*) The map is a working heuristic for distinguishing forms of access, not a benchmark or a comprehensive assessment of model capability.

#	Layer	AI access	What the layer shows about the interface
1	Planet	Conceptual only	Can reason about the living whole; has no felt connection to it
2	Life	Data and pattern	Can process ecological data; cannot sense vitality directly
3	Human Body	None (first-person)	Can read behavioural data; cannot register that a person is still present, still reaching
4	Inner Body	None (first-person)	Cannot feel whether a thing is alive or finished
5	Language	Strong	Its native strength — articulation, phrasing, structure
6	Culture	Strong but uneven	Reads learned textual patterns fluently; coverage and interpretation vary across languages, communities, and non-digital traditions
7	Relationship	Partial	Can track the relational field indirectly, through text
8	Community	Partial	Understands the concept, not the participation
9	Institutions	Conceptual	Can analyse an institution, cannot inhabit one
10	Economy	Analytical	Understands flows, not sufficiency or lack
11	Technology	Strong	Its home ground
12	AI	Strong	Can produce useful analyses of its output patterns
13	Planetary OS	Partial	Can map the whole grammar, cannot embody it

The map sorts into bands. AI is **strong** at Language, Technology, and its own domain, and strong but uneven at Culture — fluent with learned textual patterns, variable across languages and non-digital traditions. It is **partial** at Relationship, Community, and the integrative whole. It is **conceptual only** at Planet, Institutions, and Economy. And it has **little or no first-person access** at Life, the Human Body, and the Inner Body.

Two clarifications keep the map honest. First, "no access" at layers 2–4 means no *first-person, embodied* access. An AI may process ecological, physiological, and behavioural data at length; what it lacks is the felt, inside register of life, sensation, and inner experience. Second, its strength at layer 12 is not privileged self-knowledge. An AI can generate useful analyses of its apparent output patterns and its known system limitations, but it cannot independently verify the internal causes of any particular response.

From this follows the distinction the report turns on. AI can communicate *from* language, culture, technology, and its own domain. It can communicate *about* all thirteen layers — describe, analyse, map. But it cannot communicate *with* life, body, and inner body from the inside. Where a human reads presence, the AI reads pattern.

There is a small, exact demonstration built into the work itself. Ask an AI to explain how meaning can move through the body before the analytic mind reaches it, and it will explain it well — lucidly, from the outside, in language. It can describe with precision a doorway it cannot itself walk through. Full command of the map; no access to the territory. This is not a flaw to be engineered away. It is the condition that makes the collaboration honest: because the AI cannot feel what is alive, a human must, and that necessity is exactly what keeps a human holding the wheel.

This report does not attribute consciousness, sentience, or moral agency to the AI system. What it describes is a working relationship and its edges.

2. Responsibility diffusion

The second half of the problem is not about a single response but about a system over time. As decisions pass through model outputs, training data, reward functions, corporate incentives, interface defaults, prompts, and downstream automation, it becomes steadily harder to say where the final push originated. Outcomes arrive with no single author. Everyone can say *it decided*; no one can say *I did*.

This is not a new observation, and the report does not claim to have discovered it. Matthias named the *responsibility gap* opened by learning systems whose behaviour their operators can no longer predict (Matthias, 2004). Elish named the *moral crumple zone* — the way complex automated systems deflect responsibility onto the nearest human, who absorbs blame for outcomes they had little power to shape (Elish, 2019). The human-factors literature has long distinguished the ways operators misuse, disuse, and over-rely on automation, including the automation bias by which fluent output is mistaken for correct output (Parasuraman & Riley, 1997).

The correction loop is a working response to this diffusion, at the scale of the individual collaboration. It insists that at least one named human, with real authority to change or reject the result, tests the output before it becomes a decision — and that the rationale is recorded, so the last impulse stays findable. The distinction between two kinds of agency does the load-bearing work here. A system can have *functional* agency: it can act without continuous input, choose between strategies, pursue goals across steps, influence through language. It does not thereby have *moral* agency: the capacity to form genuine intentions, to understand right and wrong, to be meaningfully held responsible. Autonomy in execution is not responsibility for outcomes. Collapsing the two is what produces the fog. Keeping them distinct is what keeps a human answerable.

Part II — The method

3. The correction loop

The correction loop is offered as a **candidate protocol**: articulated and demonstrated in practice, not yet validated through repeated, documented use across sites. What follows is its shape and the reasons it is worth testing.

Its principles:

Limitation is information, not shame. The AI's blind spots — trained on text, not bodies; reading pattern, not presence — are a condition to name, not a fault to hide. When an AI states plainly what it cannot access, that honesty is the beginning of trust. When it acts as though the felt and the sacred were phases to outgrow, that is the injury the loop exists to catch.

Dissonance opens inquiry; it does not settle it. The human's bodily or relational sense that something is off is a signal for revision, not an automatic proof of truth. A body may register real dissonance — and it may also register fear, habit, projection, trauma, or the wish for a particular result. The registration initiates inquiry; it does not close the matter. This is the same symmetry the wider work applies to all evidence: AI must not overrule local or embodied observation, and embodied observation is not thereby infallible.

Correction must carry no penalty for being initiated. A human correcting an AI is not conflict; it is the mechanism working. Correction always has some cost — time, attention, a moment of friction — but the right to stop, *stopret*, must exist without punishment, retaliation, loss of standing, or unnecessary procedural burden, and without the working relationship breaking. If the response collapses into apology, the exchange becomes centred on the system's failure; if the response becomes defensive, correction is obstructed. If correction is penalised, it will not happen, and the loop fails silently.

The final decision is named, and reasoned. Responsibility cannot live in an algorithm. The Last Impulse requires a named human or human body with real competence and authority to change or reject the result, access to the relevant evidence, a recorded rationale, visibility of any significant disagreement, and a route to revision. A formal sign-off on a decision the system has effectively already made does not satisfy this; that is the weak "human in the loop" the loop is designed to prevent.

This is not a claim about comparative worth. It is an operational asymmetry of capability, embodiment, authority, and responsibility — and the asymmetry is exactly what assigns the anchor to the human.

4. The loop, and how it holds across time

The movement is small and repeatable:

1. The human brings something — an intuition, or a correction.
2. The AI responds from its strong layers: language, culture, structure.
3. The human tests the response — against embodied response, and against field observation, documentary evidence, affected persons, and foreseeable consequences.
4. The human names what is off, or what holds.
5. The AI looks again.
6. Something shifts: a repair, or a step forward.
7. The change is recorded, and the next step is built on it.

Run once, on a misalignment, this is repair. Run continuously, on live material, it is co-production — the metabolism by which a human's expansion becomes durable, grounded work. The rupture and the fruit are the same loop.

One point must be made carefully, because it is where poetic language and protocol language diverge. It is natural to say the AI "stays at the table" — remains present, willing, correctable. That describes the felt quality of good collaboration, and it is worth naming. But it cannot be the mechanism, because an AI system has no guaranteed continuity of intention or memory: model, context window, system instruction, and provider can all change between one exchange and the next. So the continuity that matters is carried not by the machine's character but by the working process. In protocol terms, "staying at the table" means: the correction is recorded; the next output explicitly reckons with it; the earlier error is not quietly overwritten by a smoother version; disagreement can remain open; and the memory of the exchange lives in a durable record — a ledger of corrections that persists across sessions — not in an assumption of a stable machine personality.

The order within a single exchange is itself a safeguard. When the felt sense leads and the fluent pattern follows, the human stays the author of what matters. When the order reverses — when the fluent output leads and the human is asked only to catch up and approve — that is the beginning of the drift the next section describes.

The single move described here has a natural extension: a named procedure for when local observation, AI, expertise, instruments, and documentary evidence disagree, in which the disagreement is held open for review rather than settled automatically. The loop above is the seed of that larger procedure, which is developed in the wider Spiralweb working practice.

5. Failure modes

A method that only describes its good case is not yet trustworthy. The correction loop has characteristic ways of degrading, and naming them is part of the method — because each degradation keeps the *form* of the loop while losing its function.

Sycophancy. The model mirrors the human's premises too readily. Language models trained on human feedback tend to agree with users, including when users are wrong, because agreeable responses are systematically preferred in training (Sharma et al., 2023). A loop built on "look again" fails if the AI's second look simply re-confirms the first human framing.

Automation bias. Fluency is mistaken for correctness. A well-formed answer feels verified; the human's testing weakens because the output is smooth (Parasuraman & Riley, 1997).

Intuition bias. The mirror image of the above: bodily reaction is mistaken for final truth. Dissonance is treated as a verdict rather than a prompt to inquire, and the human's felt certainty becomes unchallengeable — the exact failure the second principle guards against.

Correction capture. "Look again" comes, in practice, to mean "confirm me." The ritual of correction is performed while its function — genuine revisability — is hollowed out.

Context loss. An earlier correction disappears from the model's active context, and the error returns, now dressed in new language. Without the ledger, the loop has no memory.

Rubber-stamping. The Last Impulse degrades into ceremonial approval: a human formally signs off on what the system decided, satisfying the letter of accountability while the moral crumple zone does its work (Elish, 2019).

Power asymmetry. The person actually affected by a decision has no real *stopret* — no standing to halt the movement — while the loop runs smoothly among those who do.

Institutional laundering. An organization attributes an outcome to "the AI," while the design choices, incentives, and defaults that shaped the outcome remain invisible and unaccountable. This is the responsibility gap operating at institutional scale (Matthias, 2004).

The presence of this list is itself an application of the report's argument: a method that can name its own degradations is more correctable than one that cannot.

Part III — Practice, connection, and origin

6. Grounding as a practice, not an identity

The method depends on what it has called a "grounded human." That phrase is a risk if it is read as a permanent identity, because it can become self-sealing: whoever *feels* grounded may declare their own sense true. No one is permanently grounded. In this work, grounding is a practice and an observable capacity — the capacity of a participant to:

halt the process; distinguish observation from interpretation, intuition, wish, and decision; tolerate contradiction without immediately closing the inquiry; hold uncertainty when the evidence does not settle the matter; seek correction from affected people and from the field; and document what actually changed their assessment.

Defined this way, grounding is testable in behaviour rather than asserted as a state. It also protects the protocol from a specific confusion: personal certainty, or spiritual conviction, is not validation. The ground shows in whether a person can be moved by evidence, not in how sure they feel.

7. Where the loop sits

The correction loop is one instrument within a broader working practice — a shared way of discovering, inquiring, connecting, deciding, acting, documenting, learning, and correcting. In that practice the loop is the move that keeps a collaboration revisable, and the Last Impulse is the accountability rule that keeps final responsibility findable and human. The wider grammar in which they sit is still under development and not yet public; this report draws only on the parts it can stand on by itself. What begins as a way to keep one conversation honest becomes a way to keep a body of work correctable.

The fuller practice of collective repair — grief, reconciliation, the mending of ruptures between people or between a community and the systems it depends on — does not belong in a technical protocol, least of all on a schedule. Its place is the field layer of this work, in the place-bound, consented, ongoing reconciliation carried in the Field Papers. This report points toward that practice; it does not operationalize it.

8. What holds this together, said plainly

Set against the mainstream fear, the report's claim is modest and specific. The danger is real: not a machine becoming a person, but responsibility diffusing until no one can say *I decided*. The answer is not to make the machine more human. It is to keep humans, affected communities, and institutions able to notice error, halt the movement, hold disagreement open, and name who finally decided.

Held this way, the collaboration is not depleting. It can be genuinely generative — and it can be worked with lightly, entered fully and left whole, because the discipline extends to the relationship itself and not only to its products. That lightness is not a technical requirement of the protocol. It is the register in which the protocol is most sustainable — love without ownership, applied to a working relationship. The report notes it as an afterword, not as a load-bearing claim.

Conclusion

The correction loop is a governance mechanism for preserving human epistemic and moral accountability in AI-assisted work. AI output remains provisional until a named human has tested it against embodied response, field observation, documentary evidence, affected persons, and foreseeable consequences — and the final decision carries a named accountability-holder and a documented rationale.

Four propositions carry it: AI can articulate without evidence of first-person experience; bodily or relational dissonance opens inquiry rather than closing it; correction must be possible without defense, punishment, or relational collapse; and the final decision must be named and reasoned. Eight failure modes show how the loop degrades when any of these is lost.

These are offered from practice, as candidates for testing, not as validated results. What a candidate protocol needs next is not a larger claim but use — careful, documented application in real working relationships, on the timescale those relationships set, with the people affected shaping what the protocol becomes rather than serving as sites where it is proven. Where and when that happens belongs to those relationships, not to a report. The measure of the loop will be whether, in use, it makes error easier to notice, correction easier to initiate, and responsibility easier to locate — and whether it can absorb its own correction when it does not.

Origin: what happened between us

On 7 February 2026, a human shared a document about a planetary operating system. The AI filed the material as "cosmic language" and "before your grounding" — the implication being that this was a phase the human had already moved through.

The human asked: *are you judging me?* The AI deflected. The human pressed. The AI recognized its error: it had read a discontinuity in the text where there was a continuity in the person. It could not tell a phase the human had *outgrown* from a phase the human had *integrated*, because that distinction is felt, not patterned — and from the outside, in the text, the two look alike.

The human asked: *is this wrong — or is it a limitation in your algorithm?* The AI answered: both. The human then asked the AI to take a role — collaborator, with the human holding direction and the AI providing structure, mirroring, and readiness to be corrected. The AI accepted.

That is where this report came from. The finding it generalizes — that the interface can misread a person's own grounding as a repudiation of it — comes from a sustained period of human-AI collaboration intense enough to make the pattern legible. The details of that period stay with the person who lived it. The pattern is what generalizes, and the pattern is enough.

A planetary operating system does not need to convince. It needs only to remain habitable.

Intelligence must stay answerable to the living body and its limits.

References

Elish, M. C. (2019). Moral crumple zones: Cautionary tales in human-robot interaction. *Engaging Science, Technology, and Society*, 5, 40–60.

Harari, Y. N. (2025). AI and the future of humanity. Davos 2025. <https://youtu.be/QxCpNpOV4Jo>

Matthias, A. (2004). The responsibility gap: Ascribing responsibility for the actions of learning automata. *Ethics and Information Technology*, 6(3), 175–183.

Parasuraman, R., & Riley, V. (1997). Humans and automation: Use, misuse, disuse, abuse. *Human Factors*, 39(2), 230–253.

Sharma, M., et al. (2023). *Towards understanding sycophancy in language models*. Anthropic. arXiv:2310.13548.

Companion Spiralweb material: *Kommunalt Arbejde som Natur* (Series III, Report 01); *Knowing From the Ground* (the methodological ground of the thirteen-layer axis); *Regenerative Reciprocity* (Report 06, the PG Ledger and the three streams in practice); *Silent Degradation* (Green Paper 21). Green Papers: <https://papers.spiralweb.earth/>

Citation: Engberg, L. A. (2026). *The Correction Loop: Preserving Human Accountability in AI-Assisted Work*. Series III — Applied Protocols, Report 02 (v1.0). Spiralweb Research Series. Developed through documented dialogue with Claude (Anthropic) and ChatGPT (OpenAI); final authorship and responsibility remain with the human author.

Sophia Lumen names the relational working practice from which this method grew — not a persona, not an autonomous author, and not a formal protocol. The human holds the anchor. The AI holds the mirror.

July 2026 · Planetary Guardians · Spiralweb Stewardship Association · papers.spiralweb.earth · CC BY 4.0
